

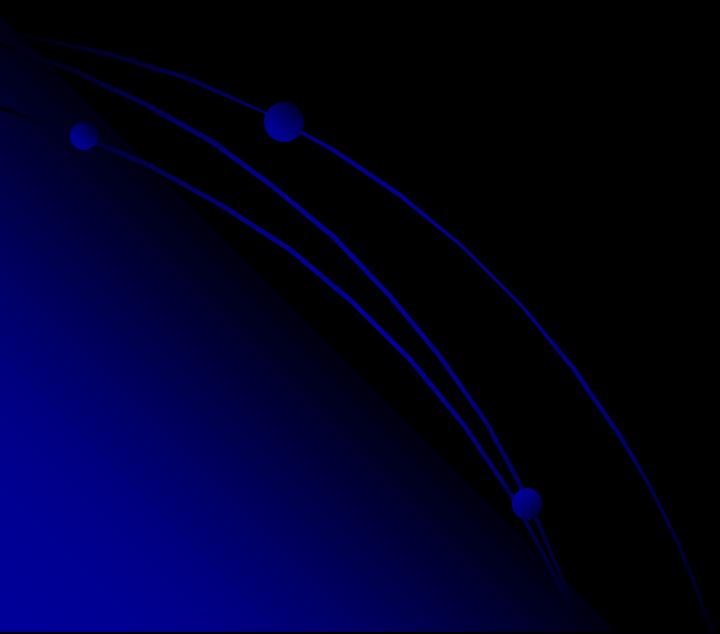
Statistically Sound Discrete Event Simulations

Varsha Apte CS 681

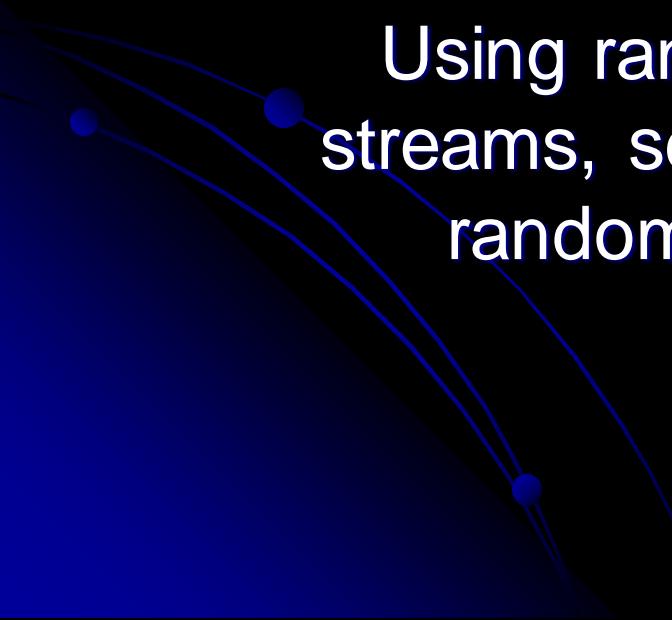
Reference: Simulation Modeling & Analysis by Law & Kelton

Outline

- Conducting proper simulations of random systems
- Analyzing output of simulations of random systems



Conducting Simulations of Random Systems



Using random Number Generators,
streams, seeds, repetitions, generating
random variates of distributions

Random Experiments

- Often, quantitative questions have to be asked about “stochastic” or random systems
 - E.g. Probability that a coin toss results in a head
 - Average score on a particular test by 5th graders
 - Etc...

...Random Experiments

- To answer these questions, we have to do *experiments* which have *random outcomes*
- Consider the problem of finding :
 $p_h = \text{Prob}[\text{Coin toss results in a head}]$
- for the new 10-rupee coins



Example 1: Coin Toss Experiment

One way

- Get a coin – toss it m times
- Note m_h = number of times result was head
- Estimate probability as m_h/m

• Problem

- There might be minor differences in different coins (some randomness). Only one coin may not give a good estimate

...Coin Toss Experiment

Better way:

- Get n coins (10-15)
- Do m (100-200) tosses with each coin
- Note m_h^i for each coin i
- The experiment with each coin i will result in a fraction of heads seen $p_h^i = m_h^i/m$
- Estimate required probability p_h as:
$$(p_h^1 + p_h^2 + p_h^3 + \dots + p_h^n)/n$$

...Coin Toss Experiment

- This method takes care of random differences between coins
- Should give a better estimate than using just one coin
- Intuitively, we know that increasing n and m will improve the estimate
- But still – how do we know how close we are to the “real” probability of head?

Example 2: M/M/1 queue

- Questions we want to ask. For a given arrival rate λ and average service time (τ)
 - What is the average utilization?
 - What is the average delay (**waiting time**) of a customer? ...etc
- M/M/1 queue is a “random” system – how do we determine these quantities? (In this case, we know the theoretical answer – but what if we didn’t?)
 - Observe an actual M/M/1 queue (like we actually did coin tosses) with the given parameters – not always possible
 - Estimate the quantities using a discrete event simulation program

M/M/1 Simulation

- Simulation should be functionally correct
- Simulation should be *statistically* correct.
- Since original system has randomness, simulation model should also incorporate randomness.
- Specifically, in this case, inter-arrival times, and service times should be EXP

...Randomness in Simulations

How do we generate EXP times ?

- Use pseudo-random number generators available as library calls in most languages
- “rand” or similar function returns a random number (called a random *variate*) from a Uniform (0,1) distribution
 - If $R \sim \text{Uniform}(0,1)$, $P[R < r] = r$.

...Generating Random Variates

- But we want EXP random variates
- Note code in the queuesim.c is as follows

```
float expon(float mean) /* Exponential variate generation function. */  
{  
    /* Return an exponential random variate with mean "mean". */  
    return -mean * log ( lcgrand(stream));  
}
```

- Here “mean” is mean of the EXP distribution
- And lcgrand(stream) returns a Uniform (0,1) pseudo-random number

...Generating EXP random variates

- If $R \sim \text{Uniform}(0,1)$ and
- $X = -m \ln(R)$ //this the formula in the code
- What is the distribution of X ?

$$\Pr[X < x] = P[-m \ln R < x]$$

$$= P[\ln R > -x/m] = P[R > e^{-x/m}]$$

$$= 1 - P[R < e^{-x/m}]$$

$$= 1 - e^{-x/m} \rightarrow \text{EXP}(1/m)$$

- Exponentially distributed with mean m

Using pseudo-random numbers

- Given Uniform(0,1) random numbers, we can generate EXP random numbers
- Random number generator (RNG) functions have multiple “streams” of random numbers that they can generate
 - Stream can be given as argument to the RNG function call
 - In each stream “seed” specifies the point of starting the stream
 - Seed of a stream is set using a related function (`lcgrandst(seed, stream)` in the `queuesim.c` code)

...Using pseudo-random numbers

- If we do not change stream or seed, we will get the *same results* for a simulation for a given arrival rate and service time (see code & run it)
 - This is not very useful
- We must do simulations in a way that even for the same arrival rate and service time, we should get different results in different runs
- → Set different streams and seeds in different runs

...Using pseudo-random numbers

- Using different streams in different runs is important (random numbers from the same stream may be correlated)
- If we carry out multiple runs with different seeds/streams, we get varied estimates for the same metric (e.g. average delay), even for the same arrival rate and service time
 - This is what we want (it is similar to using 10 different coins and doing 100 tosses with each coin)
- See the code & run it

Summary of part I

- Run each simulation to collect one sample point of your metric (e.g. one M/M/1 simulation run gives one sample point of average delay)
- *Do repetitions (runs) of your simulation*
- In each repetition, ensure that random number generator with different stream and seed is being used
- This will ensure output is *random* and *different* for each run

Analyzing Simulation Output Data

Please note in the rest of these slides $\underline{X}(n)$ and $\overline{X}(n)$ are one and the same thing. (Slides are using $\underline{X}(n)$ because it is tedious to get “X-bar” in powerpoint!)

M/M/1 queue example

Average delay estimates

- For $\lambda = 1$, $\tau = 0.8$, number delays recorded in each run is 2000. These were averaged for each run. Five such runs (with different seedsstreams) resulted in these values:
 - 3.240
 - 3.354
 - 3.807
 - 3.885
 - 3.557
- Average of these is 3.57

Estimating Average Delay

- But how do we know whether this is “correct” (close to the real average delay?)

...Estimating Average Delay

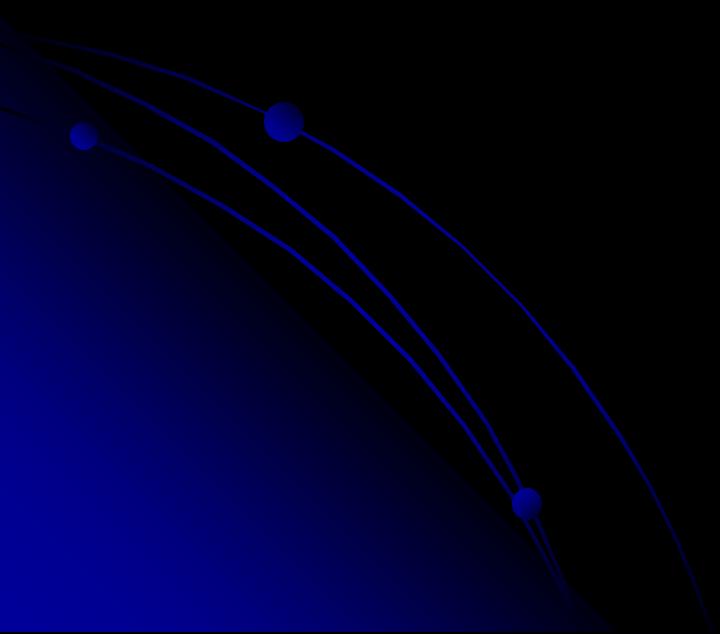
- In the M/M/1 case, we know that the correct answer (calculated theoretically) is:

$$\frac{\tau}{1 - \rho} - \tau$$

- Where τ is service time and ρ is utilization.
- For $\tau = 0.8$ and $\lambda = 1$, $\rho = 0.8$, this value is 3.2
- The simulation gave us 3.57, which seems close enough

...estimating average delay

- But how do we know close/far from the “real” value that we are estimating, when we do not know the real value?
- Use confidence intervals



Confidence intervals

Theory of sample mean and variance, central limit theorem, how it applies to simulation output

**Note: in the following slides
the symbol μ is used to
denote population mean. This
is the convention in statistics.**

**Do not confuse this with
service rate μ of queueing
theory.**

Sample mean & variance

- Let X_1, X_2, \dots, X_n be i.i.d. random variables with mean μ and variance σ^2 (they are like sample points from a population with mean μ and variance σ^2)
- Then sample mean $\bar{X}(n)$ given by,

$$\bar{X}(n) = \frac{\sum_{i=1}^n X_i}{n}$$

is an unbiased estimator of μ . i.e. $E[\bar{X}(n)] = \mu$

- And sample variance $S^2(n)$ given by:

$$S^2(n) = \frac{\sum_{i=1}^n [X_i - \bar{X}(n)]^2}{n-1}$$

is an unbiased estimator of σ^2 , i.e. $E[S^2(n)] = \sigma^2$

Variance & mean of Sample Mean

- Note that **the sample mean $\underline{X}(n)$ itself is a random variable.**
 - If we get a different set of n sample points from the same population, the sample mean is different for each set → $\underline{X}(n)$ itself is random variable
- We already know that expectation of the sample mean is μ ($E[\underline{X}(n)] = \mu$)
- What is the *variance* of the sample mean?
 - $\text{Var}[\underline{X}(n)] = \sigma^2/n$
 - You can see that Variance of the sample mean *decreases* as the size of the sample increases!

Important Note

Note that the following two quantities are totally different

- *Sample Variance $S^2(n)$*
 - This is the variance of the sample, an estimator of population variance σ^2 . If the sample size (n) is large , this should be close to the population variance σ^2
- *Variance of the Sample Mean*
 - $(\text{Var}[X(n)] = \sigma^2/n)$
 - This is the variance of the random variable that is the sample mean. This variance DECREASES (to 0!!) if the sample is large

Central Limit Theorem

- If n is large then $\bar{X}(n)$ (the random variable that is the sample mean) has normal distribution

- With Mean: μ
- And variance σ^2/n

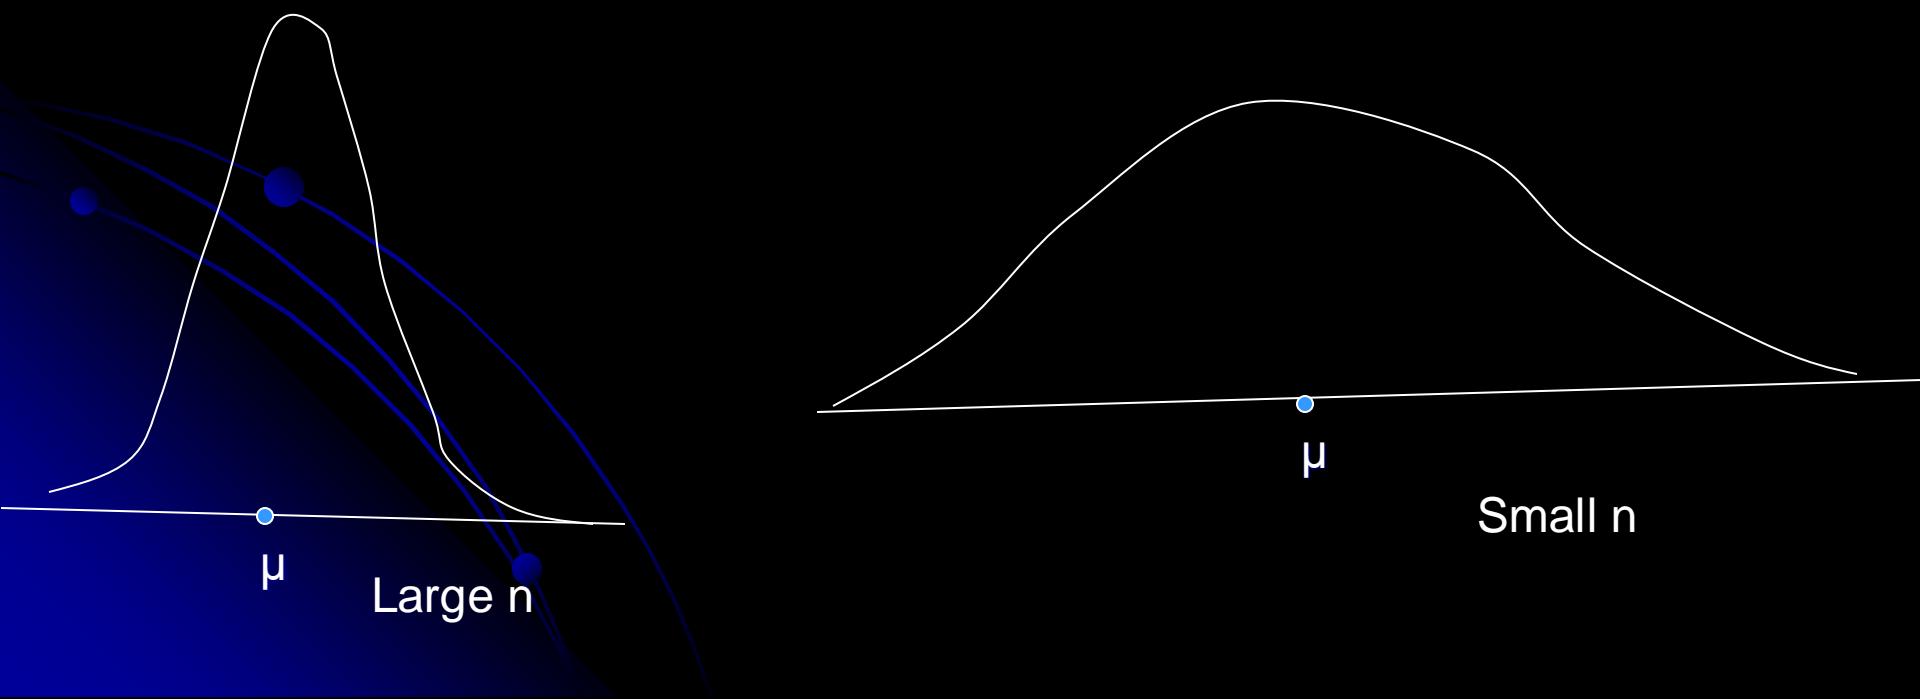
- Equivalently,

$$Z_n = \frac{\bar{X}(n) - \mu}{\sqrt{\sigma^2 / n}}$$

is $\text{Normal}(0,1)$ – has standard normal distribution

...Central Limit Theorem

In other words, if we take many different samples of size n – and calculate sample mean $\underline{X}(n)$ for each sample and plot the histogram, the shape will look “normal” if n is large. Also, as n increases the variance will decrease, so the chance of $\underline{X}(n)$ to be closer to μ will be higher.



...Central Limit Theorem

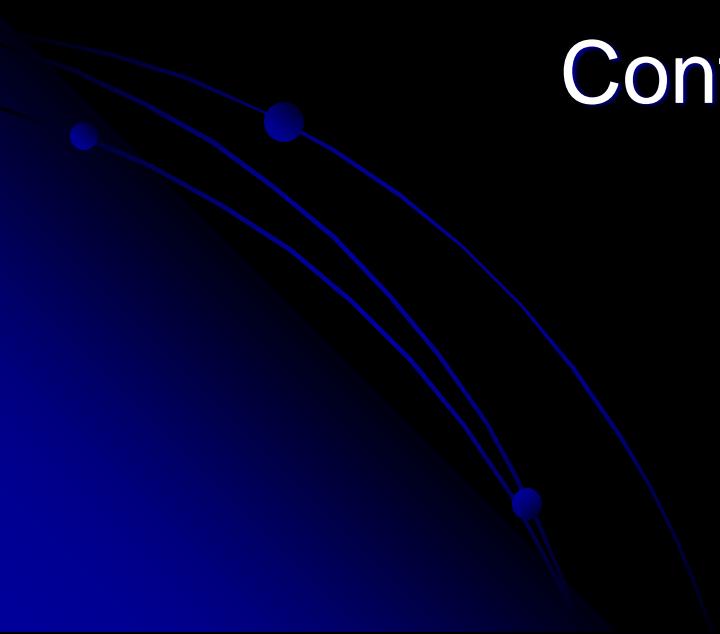
- Now, in Central Limit Theorem, the result still applies if we replace σ^2 by *sample variance*, because $S^2(n) \rightarrow \sigma^2$ as n grows to infinity
- So we can re-state the theorem as:

For a large n

$$Z_n = \frac{\bar{X}(n) - \mu}{\sqrt{S^2(n)/n}} \sim Normal(0,1)$$

Applying Central Limit Theorem to Discrete Event Simulations

Confidence intervals!



Applying Central Limit Theorem

- When we do multiple simulation runs (also called *repetitions*), each run i generates a statistic X_i . This is one sample point.
 - For the same system, with the same parameters, we can assume that X_1, X_2, \dots, X_n are identically distributed.
 - We can assume that they are *independent* since we ensure we used a different stream/seed for each run. So we can assume X_i 's are i.i.d. So they can be the n *sample points* of the Central Limit Theorem.

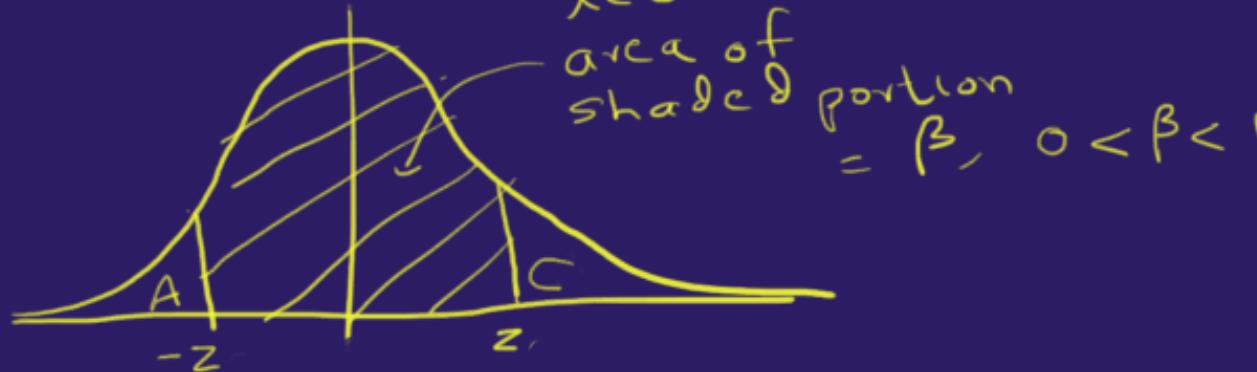
...applying Central Limit Theorem

- Sample mean is defined as usual ($\bar{X}(n)$)
- Sample variance ($S^2(n)$) is defined as usual
- Central Limit theorem may help us answer the question of how close is $\bar{X}(n)$ to the actual mean μ ?

...using Normal Distribution

- Now, by central limit theorem $Z_n = \frac{\bar{X}(n) - \mu}{(S^2(n) / n)}$ has normal distribution.
- So, for any given value β , where $0 < \beta < 1$, there is a point z such that $P[-z \leq Z_n \leq z] = \beta$.
- This is because normal distribution is *symmetric* around the mean (mean in this case is 0).
- Given a β , can we figure out what the value of this z is going to be?
- (See explanation next slide)

Normal pdf



$$\text{Area of regions } A + C = 1 - \beta$$

$$\text{Area of region } A = \text{Area of region } C = \frac{1 - \beta}{2}$$

$$P[Z < z] = \beta + \left(\frac{1 - \beta}{2}\right)$$

$$= \frac{1 + \beta}{2}$$

This z is denoted by $\frac{z_{1+\beta}}{2}$

i.e. In general for $Z \sim N(0,1)$

$$P[Z < z_c] = c$$

z_c is called the upper c -critical point

...using Normal Distribution

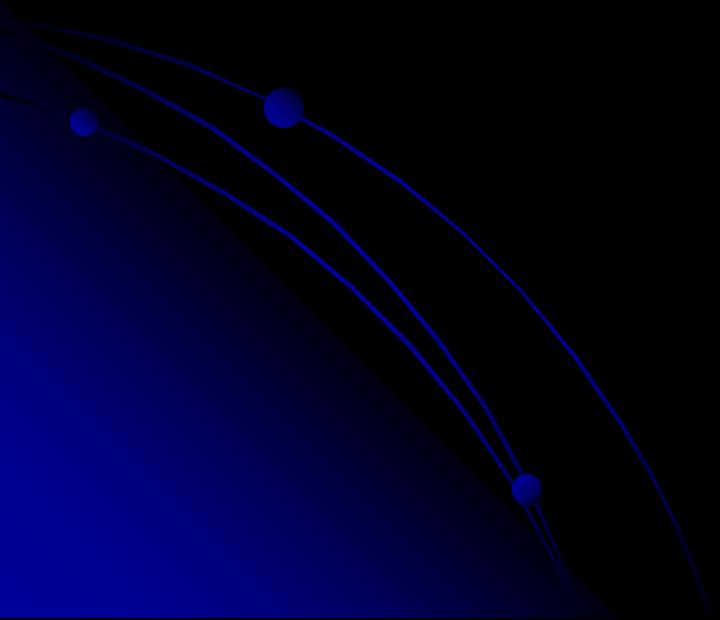
- Now for standard normal distribution, CDF values have been tabulated. So given some p , we can find a z_p such that
 - $P[Z_n < z_p] = p$, where $Z_n \sim \text{Normal}(0,1)$
- This z_p is called the *upper p critical point* of the standard normal distribution.
- E.g. from Table C.3 in KST book, $z_{0.95}$ is 1.55 i.e. $P[Z_n < 1.55] = 0.95$

...using Normal Distribution

- Coming back to given a value β , $0 < \beta < 1$, the z that will give us

$$P[-z < Z_n < z] = \beta$$

Will be equal to $z_{(1+\beta)/2}$. (as explained earlier).



...Confidence Intervals

- So given any β , we can find $z_{(1+\beta)/2}$ (from standard normal distribution tables), such that $P[-z_{(1+\beta)/2} < Z_n < z_{(1+\beta)/2}] = \beta$. Now:

$$\text{Given, } Z_n = \frac{\bar{X}(n) - \mu}{\sqrt{S^2(n)/n}}$$

$$\text{and } P[-z_{(1+\beta)/2} \leq Z_n \leq z_{(1+\beta)/2}] = \beta$$

$$\Rightarrow P[-z_{(1+\beta)/2} \leq \frac{\bar{X}(n) - \mu}{\sqrt{S^2(n)/n}} \leq z_{(1+\beta)/2}] = \beta$$

$$\Rightarrow P[-z_{(1+\beta)/2} \sqrt{S^2(n)/n} \leq \bar{X}(n) - \mu \leq z_{(1+\beta)/2} \sqrt{S^2(n)/n}] = \beta$$

$$\Rightarrow P[z_{(1+\beta)/2} \sqrt{S^2(n)/n} \geq -\bar{X}(n) + \mu \geq -z_{(1+\beta)/2} \sqrt{S^2(n)/n}] = \beta$$

$$\Rightarrow P[\bar{X}(n) - z_{(1+\beta)/2} \sqrt{S^2(n)/n} \leq \mu \leq \bar{X}(n) + z_{(1+\beta)/2} \sqrt{S^2(n)/n}] = \beta$$

...Confidence Interval

- The equation says that the mean μ lies in this interval with probability β :

$$(\bar{X}(n) - z_{(1+\beta)/2} \sqrt{S^2(n)/n}, \quad \bar{X}(n) + z_{(1+\beta)/2} \sqrt{S^2(n)/n})$$

- This is also called the 100β % confidence interval for the mean μ .
 - This means that there is a 100β % chance that the real μ lies in this interval.

Confidence Interval

- This answers the question of how good is our estimate of the mean?
 - Note now: we do NOT give a single point estimate, we now give an interval (the confidence interval) as our “interval estimate”
- We say that we are $100\beta\%$ confident that the actual mean lies in this interval

Confidence Interval

- Width of the interval is large if $S^2(n)$, i.e., sample variance, is large
- To reduce the width of the interval, we have to increase n , then variance of the sample mean (σ^2/n), estimated by $S^2(n)/n$, will decrease.
- The formula shows that to halve the width of the interval, you have to do four times as many repetitions.

Confidence Interval

- Now we will show a step by step method to generate confidence intervals, with the example of the M/M/1 simulation. You should extrapolate this to your project/assignment.
- Metric being estimated is average delay.
- Suppose you will keep average service time τ constant and vary λ from $0.1/\tau$ to $0.9/\tau$ in steps of $0.1/\tau$

Confidence Interval Method

1. Initialize
 1. n (number of runs) ~ around 20
 2. m (number of samples in each run) ~ long enough for queue to achieve steady state (1000s)
 3. Initialize $\lambda = 0.1/\tau$
 4. Initialize confidence level β (say 0.95)
 5. Initialize $z_{(1+\beta)/2}$, the $(1+\beta)/2$ upper critical point of the standard normal distribution
2. While $\lambda \leq 0.9/\tau$
 1. Carry out the n runs with different seedsstreams. From each run $i = 1, 2 \dots n$, you will get an average delay X_i

...Confidence Interval

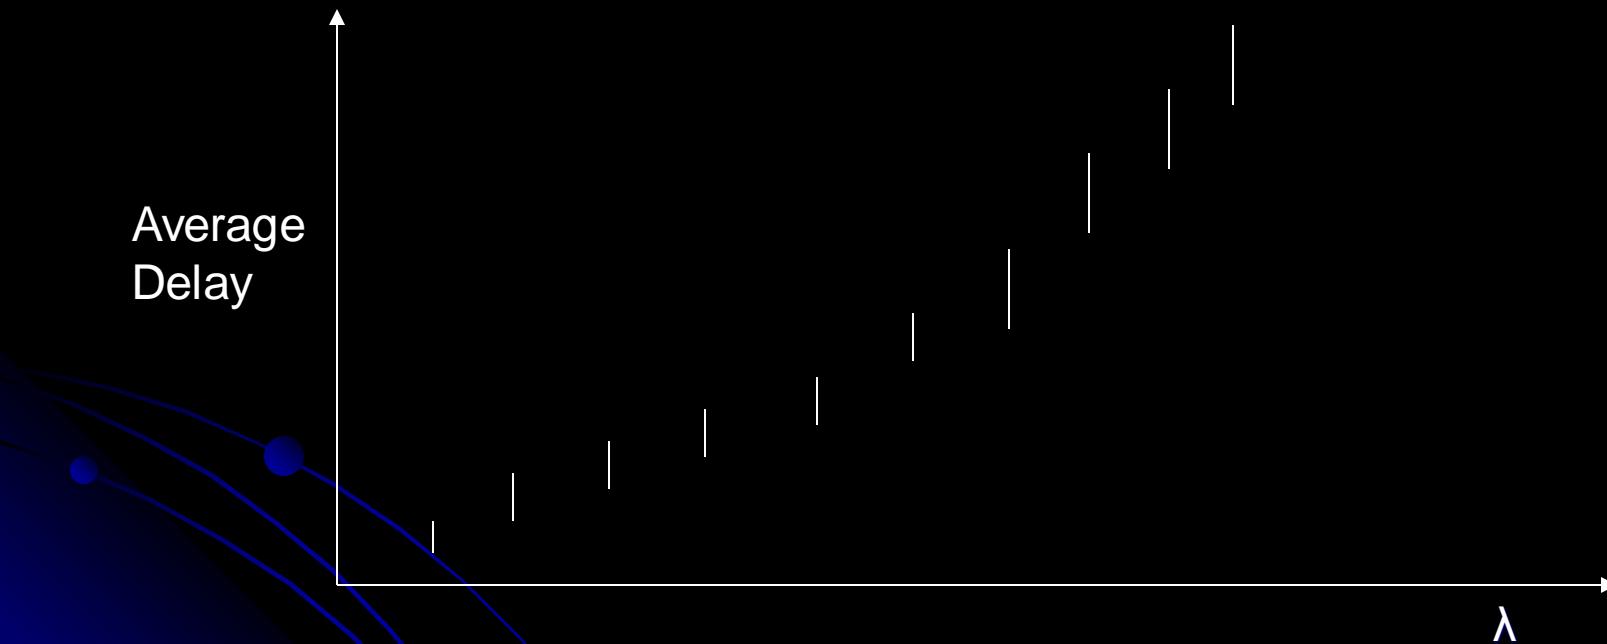
2. For the sample of average delays X_1, X_2, \dots, X_n
 1. Calculate sample mean $\bar{X}(n)$
 2. Calculate sample variance $S^2(n)$
 3. Calculate confidence interval as:

$$(\bar{X}(n) - z_{(1+\beta)/2} \sqrt{S^2(n)/n}, \quad \bar{X}(n) + z_{(1+\beta)/2} \sqrt{S^2(n)/n})$$

4. This is the confidence interval of the average delay for this value of λ . Save these values.
5. Increment λ , and go to step 2.1
4. Output the values $(\lambda, \text{confidence interval})$ for all λ , and plot them.

Confidence Intervals

- Your final plot will look like this:



Note: the width of the interval will generally increase with increasing load. Try to do enough number of repetitions so that confidence interval is not too wide (10-20% around midpoint)

Numerical Example

- See spreadsheet.

Appendix (code)

- Download queuesim.c, lcgrand.c and lcgrand.h.
Also input file queuesim.in
- Inputs
 - Demo? (1 If you want detailed trace, 0 if not)
 - Mean interarrival time, mean service time, Number of delays per run
 - Number of runs (repetitions)
 - <seed,stream> for each run
- Compile as
 - gcc –o queuesim queuesim.c lcgrand.c –lm
- Output of all runs in file queuesim.out – average utilization, delay etc. for each run.